



Secure Development On Kubernetes

Security Meetup by SBA Research
10.06.2020

Andreas Falk

meetup



Introduction

Andreas Falk
Novatec Consulting



andreas.falk@novatec-gmbh.de / [@andifalk](https://www.instagram.com/andifalk)

<https://www.novatec-gmbh.de/beratung/agile-security>

Agenda

1. What Can Go Wrong?
2. Application Security
3. Container Security
4. Kubernetes Security
5. Kubernetes Secrets

Where are the Slides and the Code?

Look here:

<https://github.com/andifalk/secure-development-on-kubernetes>

What can go wrong?

Introduction



Severe Vulnerability in Kubernetes

Source: <https://blog.aquasec.com>

Severe Privilege Escalation Vulnerability in Kubernetes



Ariel Shuper • December 06, 2018

Severe Privilege Escalation Vulnerability in Kubernetes (CVE-2018-1002105)

Earlier this week, a [severe vulnerability in Kubernetes](#) (CVE-2018-1002105) was disclosed that allows an unauthenticated user to perform privilege escalation and gain full admin privileges on a cluster. The CVE was given the high severity score of 9.8 (out of 10) and it affects all Kubernetes versions from 1.0 onwards, but fixes are available for recent versions.

Crypto Mining Via K8s Dashboard

Source: <https://blog.heptio.com>

On Securing the Kubernetes Dashboard



Joe Beda

Follow

Feb 28, 2018 · 13 min read

Recently Tesla (the car company) was alerted, by security firm RedLock, that their Kubernetes infrastructure was compromised. The attackers were using Tesla's infrastructure resources to mine cryptocurrency. This type of attack has been called "cryptojacking".

The vector of attack in this case was a Kubernetes Dashboard that was exposed to the general internet with no authentication and elevated privileges. Not only this, but core AWS API keys and secrets were visible. How do you prevent this from happening to you?

Open ETCD Ports in Kubernetes (1)

<https://shodan.io>

SHODAN etcd port:"2379" 🔍 🏠 Explore Downloads Reports

🔧 Exploits 🌐 Maps 📄 Share Search 📄 Download Results 📄 Create Report

TOTAL RESULTS
2,450

TOP COUNTRIES



China	1,116
United States	541
Germany	138
France	117
Singapore	70

TOP ORGANIZATIONS

Hangzhou Alibaba Advertisin...	417
Amazon.com	273
Tencent cloud computing	172
China Unicom Beijing	111
Hetzner Online GmbH	54

New Service: Keep track of what you have connected to the Internet. Check

47.52.241.38
Alibaba
Added on 2019-07-02 11:19:29 GMT
🇭🇰 Hong Kong
cloud

etcd
Name: etcd-hk
Version: 3.2.6
Uptime: 47h12m20.876361718s
Peers: http://10.70.10.205:2380

34.77.57.47
47.57.77.34.bc.googleusercontent.com
Halliburton Company
Added on 2019-07-02 11:05:41 GMT
🇺🇸 United States

etcd
Name: m3db_local
Version: 3.2.10
Uptime: 118h39m34.598205154s
Peers: http://0.0.0.0:2380

13.229.135.103
ec2-13-229-135-103.ap-southeast-1.compute.amazonaws.com
Amazon Data Services Singapore
Added on 2019-07-02 11:07:34 GMT
🇸🇬 Singapore, Singapore

etcd
Name: node1
Version: 3.1.0
Uptime: 20m8.52416951s
Peers: http://node1:2380

Open ETCD Ports in Kubernetes (2)




```
$ etcdctl --endpoints=http://xx.xx.xx.xx:2379  
cluster-health
```

```
member b97ee4034db41d17 is healthy: got healthy  
result  
from http://xx.xx.xx.xx:2379  
cluster is healthy
```

Vulnerable Docker Images

Source: The state of open source security report (snyk.io)



All is Root



CZnative @ home

@pczarkowski

Welcome to Kubernetes where everything runs as root and the security doesn't matter!


14:22 - 8. Mai 2019



So what can WE do as Developers?

Application- / Docker- / K8s-Security

The Path for Secure Development on K8s



The Path for Secure Development on K8s



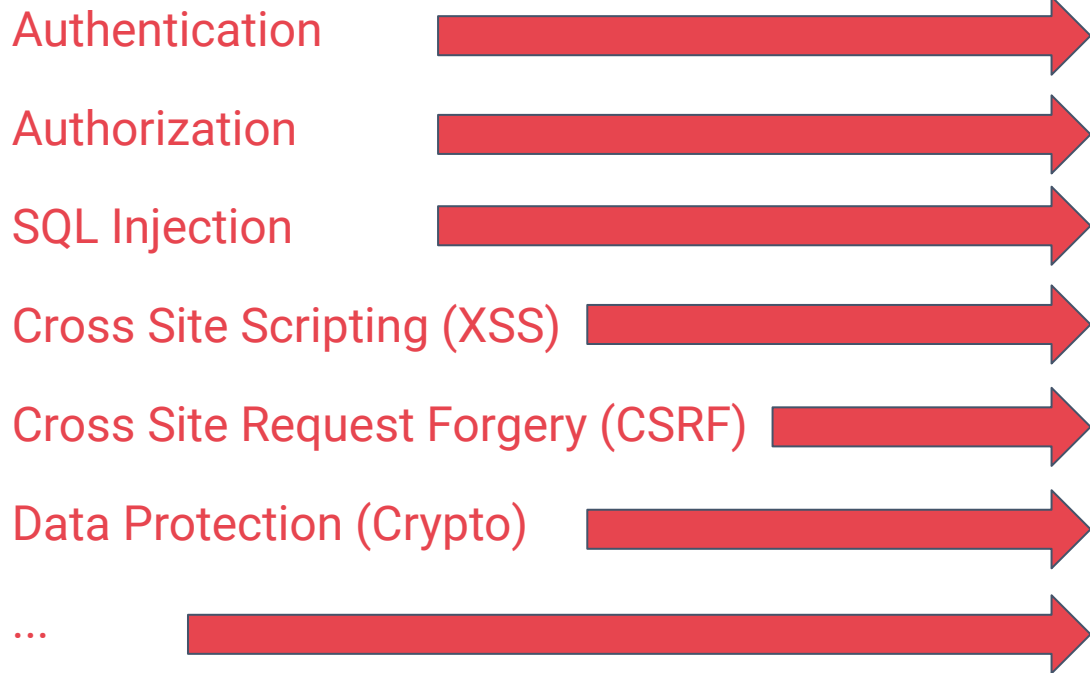
**Application
Security**

Container
Security


Kubernetes
Security

Kubernetes
Secrets

Application Security



Application Security





Live Demo: Show me the code

Iteration 1: Application Security

<https://github.com/andifalk/secure-development-on-kubernetes>

The Path for Secure Development on K8s



OWASP Docker Top 10

1. Secure User Mapping
2. Patch Management Strategy
3. Network Segmentation and Firewalling
4. Secure Defaults and Hardening
5. Maintain Security Contexts
6. Protect Secrets
7. Resource Protection
8. Container Image Integrity and Origin
9. Follow Immutable Paradigm
10. Logging

<https://github.com/OWASP/Docker-Security>

<https://doi.org/10.6028/NIST.SP.800-190>

<https://github.com/OWASP/Container-Security-Verification-Standard>

<https://www.bsi.bund.de>

NIST Special Publication 800-190

Application Container Security Guide

OWASP
Container Security Verification Standard



Bundesamt
für Sicherheit in der
Informationstechnik

SYS: IT-Systeme

SYS.1.6: Container



Virtual Machine (VM) Basics


vmware®
Workstation

KVM


vmware® ESXi



Firecracker

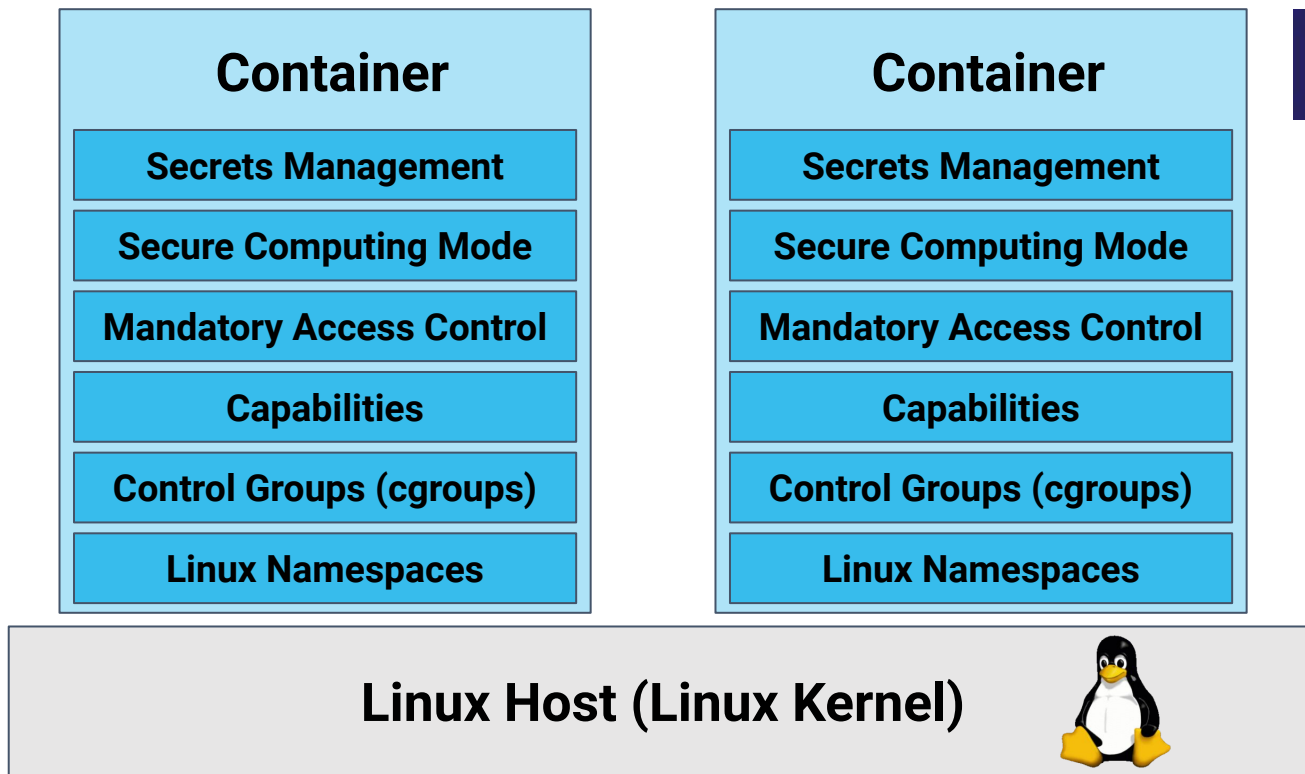


Type 1 Virtual Machine Monitor



Type 2 Virtual Machine Monitor


Container (Security) Basics



OPEN CONTAINER INITIATIVE



docker



LXD

Linux Kernel Namespaces

- Process IDs
- Network
- Mount Points
- Inter-Process Communications (IPC)
- User & Group IDs
- Unix Timesharing System (UTS): hostname & domain names
- Control groups (cgroups)

```
$ man namespaces  
$ sudo lsns
```

Linux Control Groups (cgroups)

- Resource Limits
 - CPU
 - Memory
 - Devices
 - Processes
 - Network

For Java this only works with container aware JDK versions as of **OpenJDK 8u192** or above
Recommendation: Use Java 11

Linux Capabilities

- Break up privileges into smaller units
 - CAP_SYS_ADMIN
 - CAP_NET_ADMIN
 - CAP_NET_BIND_SERVICE
 - CAP_CHOWN
 - ...

```
$ man capabilities  
$ docker run --cap-drop=ALL --cap-add=NET_BIND_SERVICE
```

<http://man7.org/linux/man-pages/man7/capabilities.7.html>

Linux Mandatory Access Control & System Calls

- Restrict System Calls
 - Secure Computation Mode (seccomp)
 - Google gVisor
- Linux Kernel Security Modules (MAC)
 - AppArmor
 - Security-Enhanced Linux (SELinux)

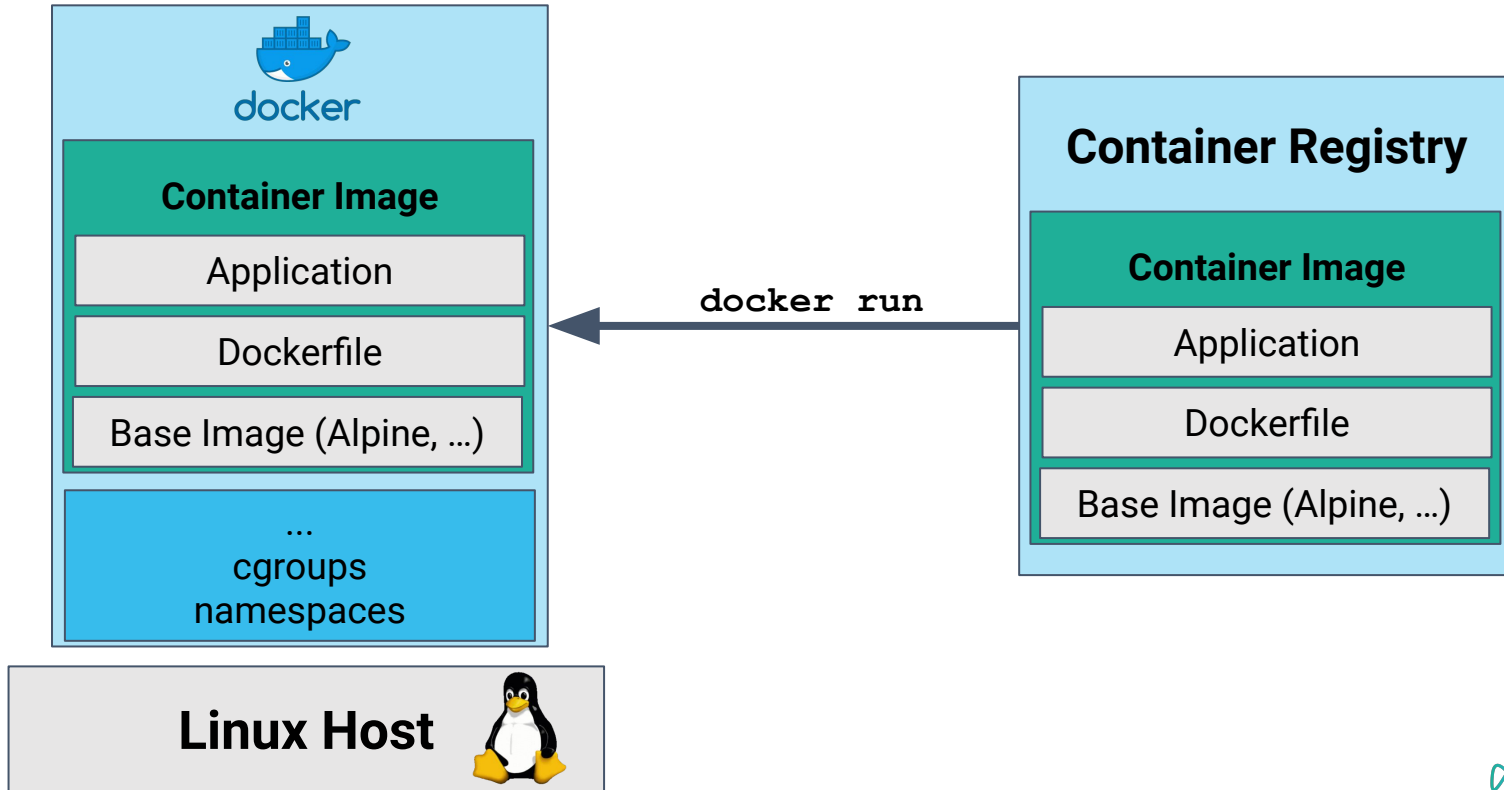
<https://docs.docker.com/engine/security/seccomp>

<https://apparmor.net>

https://en.wikipedia.org/wiki/Security-Enhanced_Linux

<https://gvisor.dev/docs>

Docker Images



Say No To Root!

USER directive in Dockerfile

```
FROM openjdk:11-jre-slim
COPY hello-spring-kubernetes-1.0.0-SNAPSHOT.jar app.jar
EXPOSE 8080
RUN addgroup --system --gid 1002 app && adduser
      --system --uid 1002 --gid 1002 appuser
USER 1002
ENTRYPOINT java -jar /app.jar
```

<https://opensource.com/article/18/3/just-say-no-root-containers>

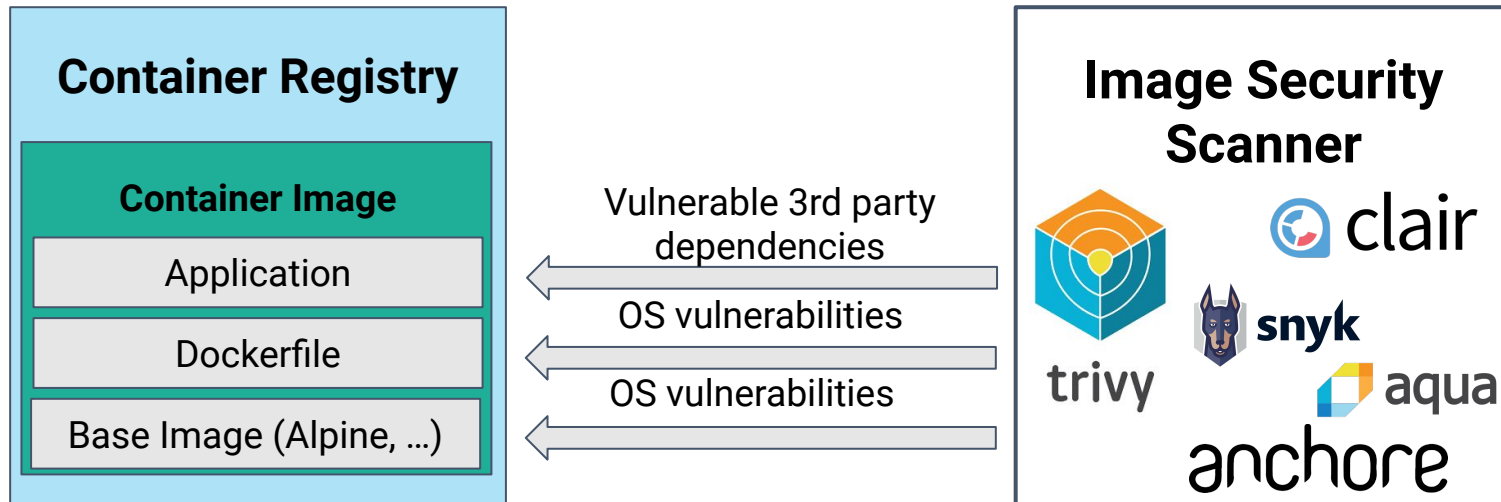
Say No To Root!

Use JIB and Distroless Images

```
plugins {  
    id 'com.google.cloud.tools.jib' version '...'  
}  
  
jib {  
    container {  
        user = 1002  
    }  
}
```

<https://github.com/GoogleContainerTools/jib>

Docker Image Security



Keep Being Secure

- Perform Image Scanning
 - Anchore
 - Clair
 - Trivy
- Use Up-To-Date Base Images

<https://anchore.com/opensource/>

<https://github.com/coreos/clair>

<https://github.com/aquasecurity/trivy>

<https://www.docker.com/blog/announcing-scanning-from-snyk-for-docker>

Docker and Snyk Announce Partnership to Streamline Container Vulnerability Scanning for Developers

Millions of developers to benefit from Snyk's vulnerability scanning natively integrated into the Docker workflow for faster and more secure application development

PALO ALTO, Calif., May 19, 2020 - [Docker](#) today announced that it has partnered with [Snyk](#) to deliver the first, native vulnerability scanning of container images in Docker. Together, Docker and Snyk will provide a streamlined workflow that makes the application development process more secure for millions of developers, allowing them to more quickly and confidently build secure applications as an automated part of their toolchain.

<https://www.docker.com/press-release/Docker-Snyk-Announce-Partnership-Container-Vulnerability-Scanning>



Live Demo: Show me the code


Iteration 2: Container Security

<https://github.com/andifalk/secure-development-on-kubernetes>

The Path for Secure Development on K8s



Kubernetes Basics



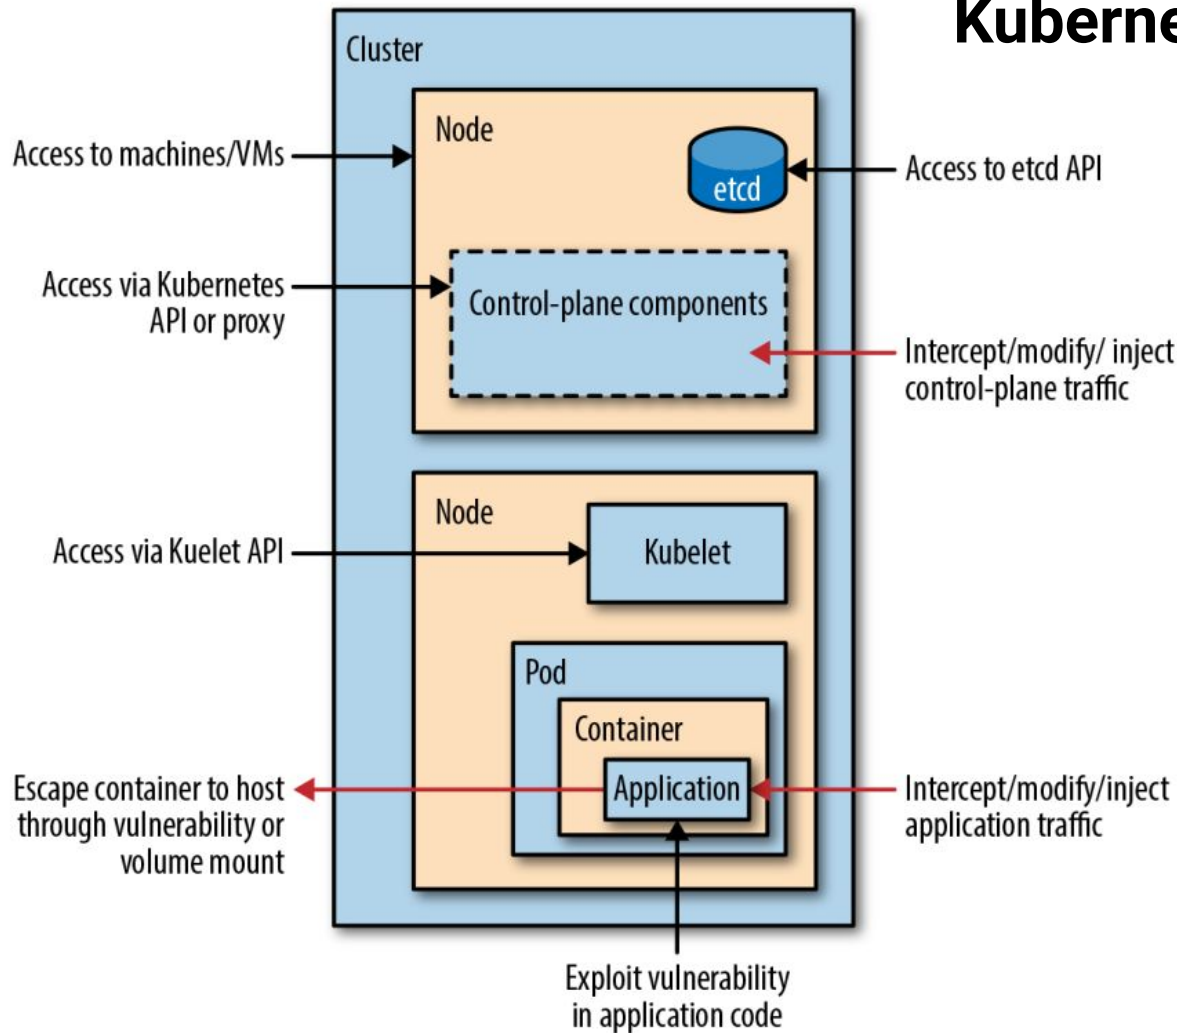
<https://kubernetes.io/docs/concepts>

<https://www.aquasec.com/wiki/display/containers/70+Best+Kubernetes+Tutorials>

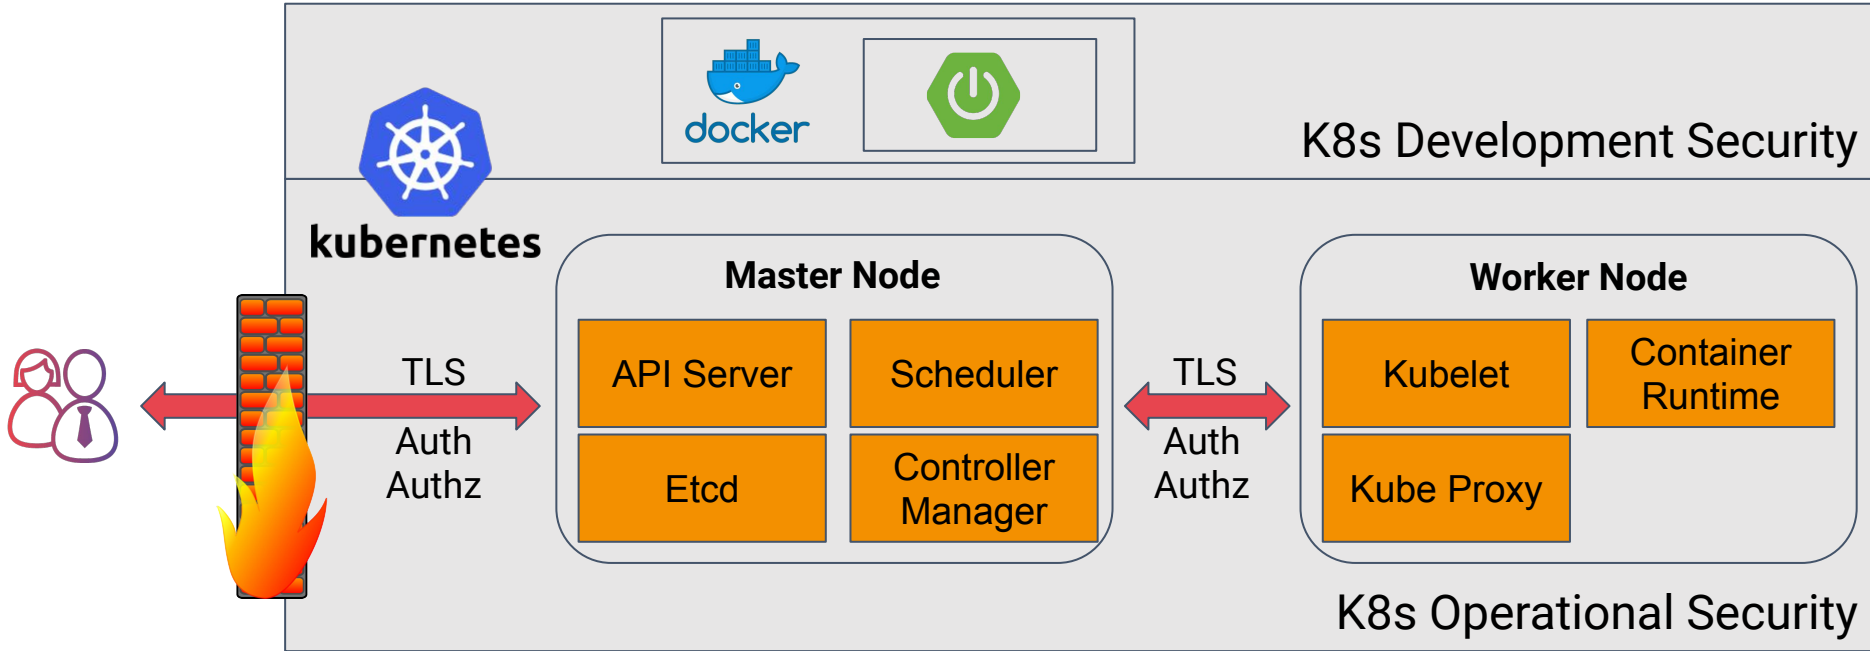
Kubernetes attack vectors

Source:

[Kubernetes Security, O'Reilly, 2018](#)




Operational / Development Kubernetes Security



<https://kubernetes.io/docs/concepts/security/overview/#the-4c-s-of-cloud-native-security>

<https://learnk8s.io/production-best-practices/>

Kubernetes Security



- ← Kubernetes Auditing
- ← Network Policies
- ← Role Based Access Control (RBAC)
- ← Resource Limits
- ← Pod Security Context
- ← Pod Security Policy
- ← Open Policy Agent

Resource Limits

```
spec:
  ...
  containers:
    resources:
      limits:
        cpu: "1"
        memory: "512Mi"
      requests:
        cpu: 500m
        memory: "256Mi"
    ...
```

<https://kubernetes.io/docs/tasks/configure-pod-container/assign-cpu-resource>

<https://kubernetes.io/docs/tasks/configure-pod-container/assign-memory-resource>

Pod/Container Security Context

```
spec:
  securityContext:
    runAsNonRoot: true
  containers:
    securityContext:
      allowPrivilegeEscalation: false
      privileged: false
      runAsNonRoot: true
      readOnlyRootFilesystem: true
      capabilities:
        drop:
          - ALL
```


<https://kubernetes.io/docs/tasks/configure-pod-container/security-context>

Pod Security Policy (Still In Beta!)

```
apiVersion: policy/v1beta1
kind: PodSecurityPolicy
metadata:
  name: no-root-policy
spec:
  privileged: false
  allowPrivilegeEscalation: false
  requiredDropCapabilities:
    - ALL
  runAsUser:
    rule: 'MustRunAsNonRoot'
  ...
```

<https://kubernetes.io/docs/concepts/policy/pod-security-policy>

Kubernetes Role Based Access Control (RBAC)



<https://kubernetes.io/docs/reference/access-authn-authz/rbac/>

Kubernetes Role Based Access Control (RBAC)

apiGroups	extensions, apps, policy, ...
resources	pods, deployments, configmaps, secrets, nodes, services, endpoints, podsecuritypolicies, ...
verbs	get, list, watch, create, update, patch, delete, use, ...

<https://kubernetes.io/docs/reference/access-authn-authz/rbac/>

Service Account

```
apiVersion: v1
kind: ServiceAccount
metadata:
  name: deploy-pod-security-policy
  namespace: default
```

<https://kubernetes.io/docs/concepts/policy/pod-security-policy/#authorizing-policies>

Pod Security Policy Role


```
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  name: no-root-policy-role
  namespace: default
rules:
  - apiGroups: ['policy']
    resources: ['podsecuritypolicies']
    verbs:     ['use']
    resourceNames:
      - no-root-policy
```

<https://kubernetes.io/docs/concepts/policy/pod-security-policy/#authorizing-policies>

Pod Security Policy Role Binding

```
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: deploy-pod-security-policy
  namespace: default
roleRef:
  kind: Role
  name: no-root-policy-role
  apiGroup: rbac.authorization.k8s.io
subjects:
- kind: ServiceAccount
  name: deploy-pod-security-policy
  namespace: default
```

Open Policy Agent



The Rego Playground

Examples ▾

```
1 package play
2
3 # Welcome to the Rego playground! Rego (
4 #
5 # Try it out:
6 #
7 # 1. Click Evaluate. Note: 'hello' is
8 # 2. Change "world" to "hello" in the
9 # 3. Change "world" to "hello" on line
10 #
11 # Features:
12 #
13 #     Examples  browse a collection
14 #     Coverage  view the policy stat
15 #     Evaluate   execute the policy w
16 #     Publish    share your playground
17 #     INPUT      edit the JSON value
18 #     (resize) DATA edit the JSON value
19 #     OUTPUT     view the result of p
20
21 default hello = false
22
23 hello {
24     m := input.message
25     m == "world"
26 }
```

INPUT

```
1 {
2   "message": "world"
3 }
```

DATA


OUTPUT

```
Found 1 result in 86.993
1 {
2   "hello": true
3 }
```

<https://www.openpolicyagent.org>

<https://play.openpolicyagent.org>

Open Policy Agent - Kubernetes Gatekeeper



Helm 3 Is Here! 😊



Ian Coldwater

@IanColdwater



Folge ich



For people who don't pay attention to the Kubernetes ecosystem: Helm 3.0 is a big deal, removing Tiller and drastically improving the security of that project. Great work, y'all!



Live Demo: Show me the code


Iteration 3: Kubernetes Security

<https://github.com/andifalk/secure-development-on-kubernetes>

The Path for Secure Development on K8s



Kubernetes Secrets



Kubernetes Secrets

```
apiVersion: v1
kind: Secret
metadata:
  name: hello-spring-cloud-kubernetes
  namespace: default
type: Opaque
data:
  user.username: dXNlcmg==
  user.password: azhzX3VzZXI=
  admin.username: YWRtaW4=
  admin.password: azhzX2FkbWlu
```

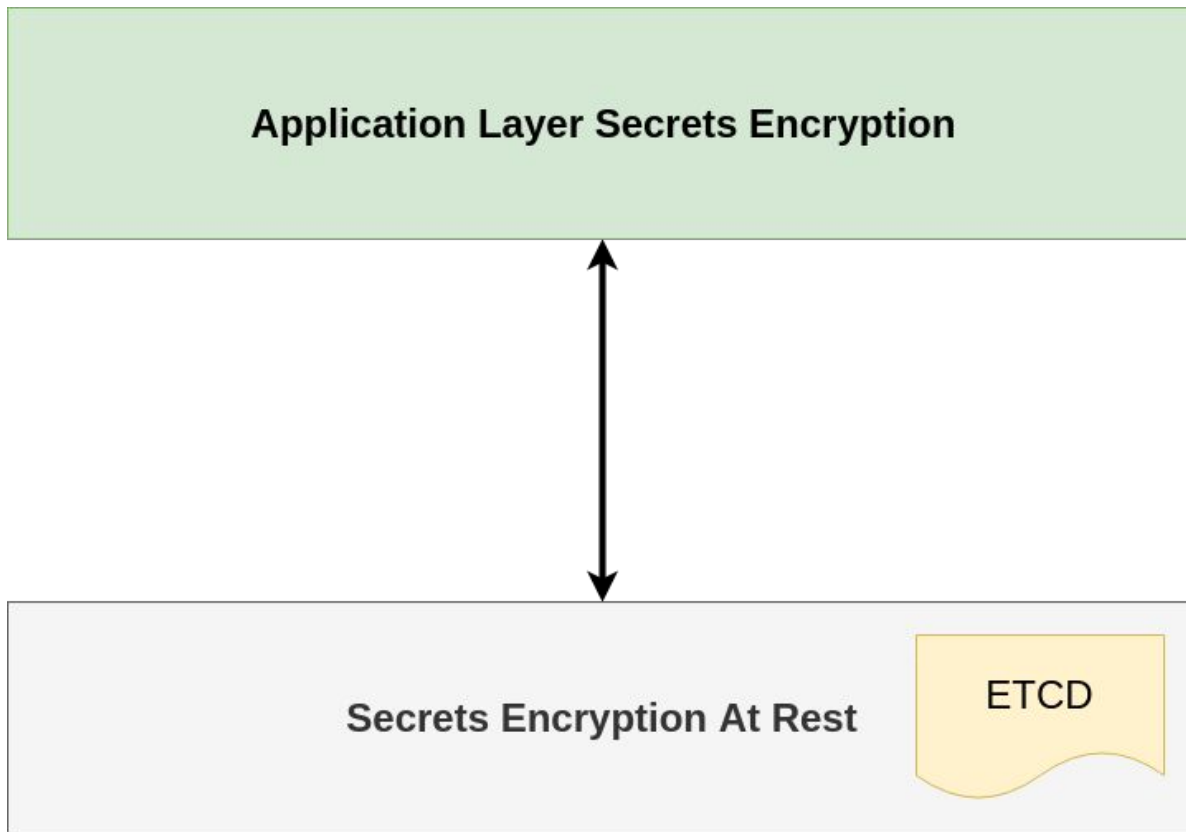
<https://kubernetes.io/docs/concepts/configuration/secret>

Kubernetes Secrets - Best Practices


- Encrypt Secret Data at Rest
Only Base64 Encoded by Default!
- Applications interacting with secrets API should be limited using RBAC
- Mount secrets instead of ENV Mapping

<https://kubernetes.io/docs/concepts/configuration/secret/#best-practices>
<https://kubernetes.io/docs/tasks/administer-cluster/encrypt-data>

Encryption Layers



Envelope Encryption On Kubernetes



<https://cloud.google.com/kms/docs/envelope-encryption>
<https://kubernetes.io/docs/tasks/administer-cluster/kms-provider>

Key Management System (KMS) Cloud Providers

- Azure Key Vault (Key Vault FlexVolume)
- Google Cloud KMS
- AWS KMS
- ...

<https://github.com/Azure/kubernetes-kms>

<https://github.com/Azure/kubernetes-keyvault-flexvol>

<https://cloud.google.com/kms>

<https://aws.amazon.com/de/kms>

What about Secrets in git

- Sealed Secrets
- Helm Secrets
- Kamus
- Sops
- Hashicorp Vault

<https://learnk8s.io/kubernetes-secrets-in-git>

<https://github.com/bitnami-labs/sealed-secrets>

<https://github.com/futuresimple/helm-secrets>

<https://github.com/Soluto/kamus>

<https://github.com/mozilla/sops>

<https://www.vaultproject.io>

Summary



Summary / Key Insights

- Containers use Linux Namespaces+Caps
- Say **NO** to root on K8s
- “**Least privilege**” for service accounts
- Keep K8s up-to-date and scan for security
- Ensure your secrets are **encrypted** in K8s
- **Scan** and keep container images **up-to-date**



Books and Online References

Books and Online References (1)

- [Kubernetes Security, O'Reilly, 2018, ISBN: 978-1-492-04600-4](#)
- [Container Security, O'Reilly, 2020, ISBN: 978-1492056706](#)
- [<https://github.com/andifalk/secure-development-on-kubernetes>](#)
- [Crafty Requests: Deep Dive Into Kubernetes CVE-2018-1002105 - Ian Coldwater \(Video\)](#)
- [Ship of Fools: Shoring Up Kubernetes Security - Ian Coldwater \(Video\)](#)
- [<https://kubernetes.io/docs/concepts/security/overview/#the-4c-s-of-cloud-native-security>](#)
- [<https://kubernetes.io/docs/tasks/administer-cluster/securing-a-cluster>](#)
- [<https://opensource.com/article/18/3/just-say-no-root-containers>](#)
- [<https://github.com/GoogleContainerTools/jib>](#)
- [<https://anchore.com/opensource/>](#)
- [<https://github.com/coreos/clair>](#)
- [<https://github.com/aquasecurity/trivy>](#)
- [\[https://www.owasp.org/index.php/OWASP_Docker_Top_10\]\(https://www.owasp.org/index.php/OWASP_Docker_Top_10\)](#)

Books and Online References (2)

- <https://kubernetes.io/docs/tasks/configure-pod-container/assign-cpu-resource>
- <https://kubernetes.io/docs/tasks/configure-pod-container/assign-memory-resource>
- <https://kubernetes.io/docs/tasks/configure-pod-container/security-context>
- <https://kubernetes.io/docs/concepts/policy/pod-security-policy>
- <https://kubernetes.io/docs/reference/access-authn-authz/rbac/>
- <https://kubernetes.io/docs/concepts/configuration/secret>
- <https://kubernetes.io/docs/tasks/administer-cluster/encrypt-data>
- <https://cloud.google.com/kms/docs/envelope-encryption>
- <https://kubernetes.io/docs/tasks/administer-cluster/kms-provider>
- <https://github.com/Azure/kubernetes-kms>
- <https://cloud.google.com/kms>
- <https://aws.amazon.com/de/kms>



Andreas Falk

Managing Consultant

Mobil: +49 151 46146778

E-Mail: andreas.falk@novatec-gmbh.de

Novatec Consulting GmbH

Dieselstraße 18/1

D-70771 Leinfelden-Echterdingen

T. +49 711 22040-700

info@novatec-gmbh.de

www.novatec-gmbh.de